MODULE 1 — LAB: HUNT THE FRONT DOOR BY HAND

> **Level coverage: L0 → L1.** Budget: 25 min. Artifacts: A1, A2, A3 (all Tier 1). Deliverables: **H-0004**, **H-0005** in your ATHF workspace.
> **Splunk time range: All time.** The telemetry is fixed at 2026-10-23 → 2026-10-30 and does not move with the calendar; a relative window (Last 24 hours, Last 7 days) returns nothing.
> All data is in `index=thrunt`. Log in with the Splunk credentials DEATHcon issued you (see the kickoff slide). Pause the video here; resume for the debrief when you're done.

SETUP (3 MIN)

Initialize your workspace if you haven't:

```
athf init
```

`athf init` asks five quick questions (workspace name, SIEM, EDR, hunt-ID prefix, retention). **Just press Enter through all of them** — the defaults (Splunk, `H-`, etc.) are exactly what this workshop expects. This creates `hunts/`, `AGENTS.md`, and the hunt template. Confirm it worked:

```
athf hunt list        # you'll see 3 seeded example hunts (H-0001–H-0003) — that's expected
```

Leave the examples; you'll add your own two (H-0004, H-0005) alongside them.

PART A — FIND THE FRONT DOOR (LO, BY HAND) (10 MIN)

You've been told developers install build tooling from a public registry. Start in the Splunk search bar. No hunt record yet — hunt like you always have.

Guiding question (orientation): *Developers install tooling all day. Before you search — what would a **malicious** install do that a normal one wouldn't? Name the behavior you're looking for, not the package.*

Step 1. Look for the install itself:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
(properties.command_line="*packmule-cli@2.4.1*" OR properties.command_line="*postinstall.js*")
| table _time hostname principal properties.parent_image_path properties.image_path properties.command_line
```

Step 2. What did the post-install script *do*? Pivot to its children:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
hostname="WS-BUILD-01"
properties.parent_image_path="*\\node.exe" properties.image_path="*\\cmd.exe"
| table _time hostname principal properties.parent_image_path properties.command_line
```

> Scoped to `WS-BUILD-01` — the build host you were told about. Staying on that one
> host now is what makes Module 2's host-agnostic pivot a genuine discovery.

Step 3. What did that shell run? The recon burst:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
hostname="WS-BUILD-01"
properties.image_path IN ("*\\whoami.exe","*\\net.exe","*\\ipconfig.exe") properties.parent_image_path="*\\cmd.exe"
| table _time hostname principal properties.parent_image_path properties.image_path properties.command_line
```

Guiding question (interpretation): *You now have three events. What is the single behavioral story that connects them — and why does it matter that it's a signed binary with a clean hash?*

PART B — THE DEAD-END (FEEL THE LO PAIN) (2 MIN)

Guiding question (reflection): *Without scrolling up — answer these honestly. If you can't, that's the lesson.*

- What were the exact three queries you just ran?
- Which hosts did you rule *out*?
- If a teammate sat down right now, could they reproduce your path?

PART C — WRITE IT DOWN (L1) (8 MIN)

Guiding question (decision): *You found one confirmed chain and one loose thread. Which becomes a "confirmed" hunt and which becomes an "open" one — and why is writing down the open one just as important?*

Convert what you found into two durable hunts. Two ways to do each — pick one, don't retype the flags:

- **Flags:** run the prepared command file, e.g. `bash commands/m1-hunt-new-H0004.sh`.
- **Template:** `cp templates/H-0004.md hunts/` — metadata is pre-filled.

> **About the numbers:** `athf hunt new` assigns the next free id. From a fresh `athf init` (three seeded examples, H-0001–H-0003) your hunts land as **H-0004 … H-0009**, and that's how every lab refers to them. If yours came out differently, use your number wherever the lab says H-000x — `athf hunt list` shows it.

H-0004 — post-install execution chain. For reference, the flags the command file runs:

```
athf hunt new \
  --title "Trojanized build-CLI post-install spawns shell + recon" \
  --technique T1195.002 \
  --tactic execution --tactic initial-access \
  --platform windows \
  --data-source edr \
  --actor "CRON GOBLIN" \
  --behavior "package-manager post-install script spawns an interactive shell that runs host/domain recon" \
  --location "developer workstations (WS-BUILD-*, WS-DEV-*)" \
  --evidence "Sysmon/eCAR process create: node.exe → cmd.exe → whoami/net/ipconfig" \
  --threat-context "packmule-cli@2.4.1 backdoored via compromised publish token; signed binary, clean hash, no CVE" \
  --hunter "<your-name>" \
  --non-interactive
```

Open the created file in `hunts/` and fill the **Check** and **Keep** sections: paste the three SPL queries, the hosts you confirmed, and the verdict.

> **On "verdict":** it isn't a literal field name — it's the record's **Status** (`planning` → `completed`) plus the one-line finding in the **Keep / Findings** section (`confirmed` / `not-found` / `needs-follow-up`).

H-0005 — unusual outbound after install. Before you write it, look at it once (still Tier 1):

```
index=thrunt sourcetype="apache:access:json" source="*proxy_access.log"
server="cdn-packmule-telemetry.net"
| timechart span=5m count by client
```

Don't solve it — that's Module 3 territory. Capture the *observation* as a hunt so it's on record (`bash commands/m1-hunt-new-H0005.sh` or `cp templates/H-0005.md hunts/`), and mark H-0005's verdict **needs-follow-up**. For reference:

```
athf hunt new \
  --title "Periodic outbound to CDN-lookalike after build-CLI install" \
  --technique T1071.001 \
  --tactic command-and-control \
  --platform windows \
  --data-source proxy \
  --actor "CRON GOBLIN" \
  --behavior "regular-interval HTTPS to a young CDN-lookalike domain following post-install execution" \
  --location "WS-BUILD-01 egress via PROXY-01" \
  --evidence "proxy access log: ~5-min-interval requests to cdn-packmule-telemetry.net" \
  --hunter "<your-name>" \
  --non-interactive
```

PART D — VERIFY (2 MIN)

```
athf hunt list          # H-0004 and H-0005 should both appear (plus the 3 seeded examples)
athf hunt validate      # every hunt should pass structural validation
```

You should leave Module 1 with **H-0004 and H-0005 on disk** — one confirmed, one open. Resume the video for the debrief.

THE AGENTIC MOMENT (RECAP)

Guiding question (reflection): *You found A1–A3 by hand at L0. What specifically was lost the moment you'd have closed your laptop — and what did* `athf hunt new` *actually change?*

Manual dead-end: you found A1–A3 by hand, but your queries lived in tabs and your reasoning lived in your head — unreproducible, unshareable, gone when you close the laptop.

Agentic recovery (L1): `athf hunt new` turned the hunt into a durable record. It's not AI yet — it's the *substrate* AI needs.

MODULE 1 — BONUS / STRETCH HUNTS

> For hunters who clear H-0004 and H-0005 early. All still Tier 1–2, still in `index=thrunt`. See `seeding/TIER_MANIFEST.md` for the full answer key.

BONUS 1 — ENUMERATE EVERY INFECTED HOST (TIER 2 PIVOT)

Guiding question (decision): *You confirmed WS-BUILD-01. Was it the only one? Decide how to look for more victims WITHOUT anchoring on the host or package you already have.*

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
properties.parent_image_path="*\\node.exe" (properties.image_path="*\\cmd.exe" OR properties.image_path="*\\sh.exe")
| stats count values(hostname) as hosts by principal
```

BONUS 2 — BASELINE THE NOISE (BEHAVIORAL RIGOR)

Guiding question (interpretation): *How often does* `node.exe` ***legitimately*** *spawn a shell in this environment? Run the pivot without the malicious filter — what does normal look like, and how would you defend "this is weird"?*

Baseline the node→shell rate across the whole environment and window — the same pivot as Bonus 1, minus any malicious anchor:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
properties.parent_image_path="*\\node.exe" (properties.image_path="*\\cmd.exe" OR properties.image_path="*\\sh.exe" OR properties.image_path="*\\powershell.exe")
earliest="10/23/2026:00:00:00" latest="10/31/2026:00:00:00"
| bucket _time span=1d
| stats count as spawns dc(hostname) as hosts by _time
| sort 0 _time
```

BONUS 3 — WRITE THE RULED-OUT HUNT

Guiding question (decision): *Find one thing that looked suspicious but was benign. Why is writing up a dead-end a real deliverable and not wasted effort?*

MODULE 2 — LAB: REUSE THE BEHAVIOR, FIND THE HIDDEN VICTIM

Level coverage: L2. Budget: 20 min. Artifact: A4 (Tier 2). Deliverable: **H-0006**.
Splunk time range: All time. The telemetry is fixed at 2026-10-23 → 2026-10-30 and does not move with the calendar; a relative window (Last 24 hours, Last 7 days) returns nothing.
Requires your Module 1 hunts (H-0004, H-0005) on disk. The A5 persistence stretch lives in `BONUS.md` and comes back as Module 3 Part B.

PART A — ORIENT OFF YOUR OWN CORPUS (L2) (4 MIN)

Guiding question (orientation): *Before you query anything new — what has your own repo already told you, and which parts of the intrusion have you provably NOT looked at yet?* The point of L2 is that your last hunt tells you where to go next.

Ask your hunt repo what you've already done:

```
athf similar --hunt H-0004
athf hunt coverage --detailed
```

Now the free-text form, to feel how the corpus answers a behavioral question:

```
athf similar "post-install shell spawns recon child process"
```

PART B — THE DEAD-END (WHY KEYWORD SEARCH FAILS) (2 MIN)

Guiding question (decision): *You want to find other victims. Your instinct is to search the indicators you already have (host, package, date). Predict what that returns — and decide whether it can possibly teach you anything new.*

Try it the naive way — reuse the *indicators* from H-0004, with the time picker on **All time**:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
hostname="WS-BUILD-01" properties.command_line="*packmule-cli@2.4.1*"
```

PART C — PIVOT ON BEHAVIOR, NOT INDICATORS (L2) (7 MIN)

Guiding question (orientation): *Take the shape you documented in H-0004 — trusted-dev-process → unexpected shell → recon. Before you run it host-agnostic and time-agnostic across the whole environment, what do you EXPECT to fall out that the keyword search couldn't reach?*

No hostname, no date, no package name:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
properties.parent_image_path="*\\node.exe" (properties.image_path="*\\cmd.exe" OR properties.image_path="*\\sh.exe")
| table _time hostname principal properties.image_path properties.command_line
```

Now take each `cmd.exe` you find and look at *its* children across all hosts:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
properties.parent_image_path="*\\cmd.exe" properties.image_path IN ("*\\whoami.exe","*\\systeminfo.exe","*\\net.exe","*\\ipconfig.exe")
| stats min(_time) as first values(properties.image_path) as cmds by hostname principal
| eval first=strftime(first,"%Y-%m-%d %H:%M:%S")
```

> The `strftime` renders `first` as a readable date instead of raw epoch — so the
> ~25-hour gap between the two hosts is something you can see at a glance.

Guiding question (interpretation): *A second host appears. What specifically makes it the same intrusion and not a coincidence — and what did you have zero indicators for a moment ago?*

PART D — WRITE H-0006 (5 MIN)

Guiding question (decision): *What is the ONE methodology note that makes this hunt reusable by the next hunter — the sentence that captures why keyword search failed and behavior won?*

`bash commands/m2-hunt-new-H0006.sh` or `cp templates/H-0006.md hunts/`, then fill in `--location` / the location field with the second host you found. For reference:

```
athf hunt new \
  --title "Second-wave post-install recon on different host (behavioral pivot)" \
  --technique T1059 \
  --tactic execution --tactic discovery \
  --platform windows \
  --data-source edr \
  --actor "CRON GOBLIN" \
  --behavior "same post-install-shell-to-recon lineage as H-0004, time-shifted onto a different developer workstation" \
  --location "<the second host + user you found in Part C> — NOT the WS-BUILD-01 host from H-0004" \
  --evidence "Sysmon node.exe → cmd.exe → whoami/systeminfo, host-agnostic pivot" \
  --threat-context "found via athf similar + behavioral pivot off H-0004, not keyword search" \
  --hunter "<your-name>" \
  --non-interactive
```

Then: in the record, explicitly write the methodology note.

PART E — VERIFY (2 MIN)

```
athf hunt list           # now 3 of your hunts
athf hunt coverage        # execution now has 2 hunts; persistence still empty
```

Persistence is still an empty coverage row. Notice it; Module 3 fills it. (Ahead of schedule? `BONUS.md` Bonus 1 lets you take the A5 persistence hunt early.) Resume the video for the debrief.

THE AGENTIC MOMENT (RECAP)

Guiding question (reflection): *Why did every indicator-anchored search miss the second victim, and what exactly did writing H-0004 down in Module 1 buy you here?*

Manual dead-end: the second victim was time-shifted and host-shifted; every keyword/time search anchored on H-0004's indicators missed it.

Agentic recovery (L2): `athf similar` + `athf hunt coverage` let you pivot on the *behavior you documented*, not the indicators you observed — and the documented behavior found a victim you had zero indicators for. Coverage also named your open gaps (persistence, C2, exfil), which is exactly what the Module 4 agent will read to propose your next move.

MODULE 2 — BONUS / STRETCH HUNTS

> For hunters who finish H-0006 early. Tier 2. See `seeding/TIER_MANIFEST.md` for the answer key.

BONUS 1 — CLOSE THE PERSISTENCE GAP EARLY (A5, TIER 2)

Guiding question (decision): `athf hunt coverage` *shows persistence empty. You have a behavior that keeps working — "unusual child of a trusted dev process." Where do you point it to close persistence, without waiting for Module 3?*

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
(properties.command_line="*core.hooksPath*") OR (properties.parent_image_path="*\\cmd.exe" properties.command_line="*loader.js*--beacon*")
| table _time hostname principal properties.parent_image_path properties.image_path properties.command_line
```

BONUS 2 — RULE OUT THE SIMILARITY FALSE-POSITIVE (RH-BENIGN-POSTINSTALL, RED HERRING)

Guiding question (interpretation): *A behavior-similarity search on H-0004's structure surfaces a second candidate. Same shape — but is it the same behavior? What single detail decides malicious vs. benign here?*

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
(properties.command_line="*postinstall*" OR properties.command_line="*node-gyp*" OR properties.command_line="*npm install*")
| table _time hostname principal properties.parent_image_path properties.image_path properties.command_line
```

BONUS 3 — QUANTIFY THE TIME-SHIFT

Guiding question (interpretation): *You found a second wave in Part C. Compute the delay between the two waves — does that interval tell you anything durable about the actor's operational cadence?*

Compute `first-seen(second host) - first-seen(WS-BUILD-01)`.

MODULE 3 — LAB: LET RESEARCH CRACK THE BEACON

> **Level coverage: L3.** Budget: 15 min. Artifacts: A5 (Tier 2, while research runs), A8 (Tier 3). Deliverable: **H-0007**.
> **Splunk time range: All time.** The telemetry is fixed at 2026-10-23 → 2026-10-30 and does not move with the calendar; a relative window (Last 24 hours, Last 7 days) returns nothing.
> Requires H-0005 (the open beacon thread) on disk. Uses your LLM provider.
> **Start Part A's research command first, then read ahead while it runs.**

PART A — GENERATE RESEARCH FROM THE OPEN THREAD (L3) (5 MIN)

Guiding question (orientation): *You have two open threads: an uncharacterized beacon (H-0005) and an empty persistence coverage row. Before you run research — what specifically don't you know how to look for yet? Point research at the gap in your own tradecraft.*

```
athf research new \
  --topic "HTTPS C2 beacon detection via TLS JA3 pivoting and shared infrastructure; git-hook persistence in developer environments" \
  --technique T1071.001 \
  --depth basic --no-web-search
```

> `--no-web-search` keeps this inside the budget. If `athf research` errors on your provider, skim the reference research doc in the workbook appendix and keep moving — don't burn the module debugging keys. Bring the error to the live Q&A.

/\/ redbeard.io Read Ahead: while the research command runs, then skim the resulting `R-XXXX` doc in `athf/research/`.

Guiding question (interpretation): *The doc will contain both real tradecraft AND plausible-but-invented specifics. How do you tell them apart — and what's your rule for when a research claim becomes a finding?*

PART B — HUNT PERSISTENCE WITH WHAT RESEARCH TOLD YOU (A5) (3 MIN)

Guiding question (decision): *Research pointed at git-hook persistence. You've hunted "unusual child of a trusted dev process" before. How does that same behavior apply to git?*

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
(properties.command_line="*core.hooksPath*") OR (properties.parent_image_path="*\\cmd.exe" properties.command_line="*loader.js*--beacon*")
| table _time hostname principal properties.parent_image_path properties.image_path properties.command_line
```

PART C — THE SHARED-INFRASTRUCTURE PIVOT (A8, TIER 3) (4 MIN)

Guiding question (decision): *Go back to H-0005's beacon. Searching the same domain again returns the same domain. Research told you actors reuse infrastructure — what do you pivot on instead of the name?*

Resolve the domain to the IP that's actually hosting it, then pivot on the neighboring addresses instead of the domain name:

```
index=thrunt sourcetype="bro:ssl:json" server_name="cdn-packmule-telemetry.net"
| stats count values(server_name) as domain by id.resp_h
```

```
index=thrunt sourcetype="bro:ssl:json"
| where cidrmatch("203.0.113.44/30", 'id.resp_h')
| stats count values(id.resp_h) as ips by server_name
```

Guiding question (interpretation): *A second domain falls out sharing the same hosting block. Why is shared infrastructure stronger evidence of "same actor" than shared domain reputation or request volume would be?*

PART D — WRITE H-0007 AND VERIFY (3 MIN)

Guiding question (decision): *This hunt came from research, not from your own data. What field makes that provenance durable, and why does linking it matter two modules from now?*

`bash commands/m3-hunt-new-H0007.sh` or `cp templates/H-0007.md hunts/`, then fill in the second domain and your `R-XXXX` id. For reference:

```
athf hunt new \
  --title "Git-hook persistence + infra-linked secondary C2 channel" \
  --technique T1546 \
  --tactic persistence --tactic command-and-control \
  --platform windows \
  --data-source edr --data-source network \
  --actor "CRON GOBLIN" \
  --behavior "pre-commit git hook relaunches beacon; secondary low-slow C2 domain shares the primary's hosting block (203.0.113.44/30)" \
  --location "WS-BUILD-01; egress to cdn-packmule-telemetry.net + <the second domain you surfaced in Part C>" \
  --evidence "Sysmon core.hooksPath→loader.js; Zeek ssl two look-alike domains on adjacent C2 IPs in the /30 you pivoted on" \
  --research "<R-XXXX from Part A>" \
  --hunter "<your-name>" \
  --non-interactive
```

```
athf hunt list              # 4 of your hunts
athf hunt coverage          # persistence + C2 now covered; cred-access + exfil still open
```

Resume the video for the debrief.

THE AGENTIC MOMENT (RECAP)

Guiding question (reflection): *The low-and-slow tasking channel (A8) is six events over days on a reputable-looking domain. Why was it invisible to manual hunting, and what exactly did the research agent contribute that you couldn't?*

Manual dead-end: the low-and-slow tasking channel (A8) is six events over days on a reputable-looking domain — invisible by volume, reputation, or keyword. You didn't know infrastructure pivoting was even the move.

Agentic recovery (L3): `athf research new` surfaced the infrastructure-reuse tradecraft (pitched as JA3/TLS-fingerprint pivoting); you adapted it to the telemetry you actually have — no `ja3` field, so you pivoted on the shared hosting IP block instead — and confirmed it in your own data. The agent taught you *what to look for*; you translated it to your environment; Splunk proved it real. And you linked the hunt back to the research (`--research`), so the reasoning chain is durable.

MODULE 3 — BONUS / STRETCH HUNTS

> For hunters who finish H-0007 early. Tier 2–3. See `seeding/TIER_MANIFEST.md`.

BONUS 1 — CHARACTERIZE THE BEACON STATISTICALLY

Guiding question (interpretation): *You confirmed periodicity by eye. What statistic would turn "looks periodic" into "is a beacon" — and what value range proves it?*

```
index=thrunt sourcetype="apache:access:json" source="*proxy_access.log"
server="cdn-packmule-telemetry.net"
| sort 0 _time | streamstats current=f last(_time) as prev by client
| eval delta=_time-prev | stats count avg(delta) p50(delta) stdev(delta) by client
```

BONUS 2 — ASK RESEARCH A BAD QUESTION ON PURPOSE

Guiding question (reflection): *Run research on a CVE that doesn't exist. What does the agent do with a false premise — and how did you catch it?*

```
athf research new --topic "does CRON GOBLIN use CVE-2024-99999 for initial access"
```

BONUS 3 — RULE OUT THE LEGIT CDN (RED HERRING RH-LEGIT-CDN)

Guiding question (decision): *Not everything hitting a CDN-looking domain is C2. Find the benign CDN traffic and state — in one line — the exact discriminator that separates it from A7/A8.*

MODULE 4 — LAB: LET THE AGENT FIND YOUR BLIND SPOTS

Level coverage: L4. Budget: 15 min. Artifacts: A9, A10 (both Tier 3).
Deliverable: **H-0008**.
Splunk time range: All time. The telemetry is fixed at 2026-10-23 → 2026-10-30 and does not move with the calendar; a relative window (Last 24 hours, Last 7 days) returns nothing.
Requires your four hunts (H-0004..H-0007) on disk — the agent reasons over them. Requires an MCP-capable client + working provider. **Do the MCP config during the kickoff setup, not now** — if the handshake fails here, use the CLI fallbacks and keep moving.

PART A — TALK TO YOUR REPO (L4) (3 MIN)

Your MCP client (Claude Code is the supported default) is pointed at `athf-mcp-server` — the config snippet is in `providers/`. The agent has these ATHF tools: `generate_hunt_hypothesis`, `search_hunts_by_text`, `find_similar_hunts`, `list_coverage_gaps`, `propose_next_query`.

Guiding question (orientation): *You're about to ask your repo a question in plain English instead of calling a tool by name. What do you expect the agent to do with "What have I hunted before?" — and what does it mean that you don't pick the tool?*

```
What have I hunted before?
```

> **Fallback (handshake failed / NL misfires):** `athf similar "beacon"` and `athf hunt coverage --detailed` CLI-side — same data, no agent. Bring the MCP error to the live Q&A.

PART B — LET THE AGENT NAME THE GAP (L4) (7 MIN)

Guiding question (decision): *You've found initial access, execution, persistence, C2. Are you done? Notice WHY this is the hardest kind of gap to see — nothing points at what's missing. Then ask your repo the question every hunter should ask and almost never can.*

```
Based on what I've already hunted, what should I hunt next?
```

Then push further, still in plain English:

> Give me a starting query for the credential-access and exfiltration gaps.

> **Fallback (NL misfires / offline):** `Use list_coverage_gaps` and `Use propose_next_query focused on credential-access and exfiltration`.
> CLI: `athf hunt coverage --detailed`.

Take the proposed direction and hunt the two Tier-3 artifacts.

Credential access (A9):

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
(properties.command_line="*type*\\.env*" OR properties.command_line="*\\.aws\\credentials*" OR properties.command_line="*findstr*password*.netrc*")
| table _time hostname principal properties.parent_image_path properties.image_path properties.command_line
```

DNS-tunnel exfil (A10) — the hardest artifact of the campaign:

```
index=thrunt sourcetype="bro:dns:json" qtype_name=TXT query="*.exfil-packmule.net"
| eval label=mvindex(split(query,"."),0) | eval len=len(label)
| stats count avg(len) as avg_label_len by id.orig_h
```

Guiding question (interpretation): *These two hits had no open thread pointing at them. Now that you see them, what makes A10 (DNS-tunnel exfil) the hardest artifact of the whole campaign?*

PART C — WRITE H-0008 (3 MIN)

`bash commands/m4-hunt-new-H0008.sh` or `cp templates/H-0008.md hunts/`. For reference:

```
athf hunt new \
  --title "Credential access + DNS-tunnel exfil (coverage-gap driven)" \
  --technique T1552.001 \
  --tactic credential-access --tactic exfiltration \
  --platform windows \
  --data-source edr --data-source network \
  --actor "CRON GOBLIN" \
  --behavior "beacon reads .env/.aws/.netrc, then exfiltrates via encoded TXT DNS queries to a lookalike domain" \
  --location "WS-BUILD-01; DNS egress to exfil-packmule.net" \
  --evidence "Sysmon cred-file reads; Zeek dns high-entropy TXT subdomains" \
  --threat-context "found via MCP list_coverage_gaps naming uncovered phases + propose_next_query" \
  --hunter "<your-name>" \
  --non-interactive
```

PART D — VERIFY (2 MIN)

```
athf hunt list          # 5 of your hunts
athf hunt coverage       # ALL SIX phases now covered
```

You've now reconstructed the entire first campaign — six kill-chain phases, five documented hunts. **Keep this workspace open.** In Module 5, you pivot to a second campaign from the same actor — and this repo becomes your head start. Resume the video for the debrief.

> **Ahead of schedule?** Ask your repo something off-script — *"Which of my hunts share infrastructure?"*, *"Summarize this whole intrusion in three sentences for my IR lead."* That's `BONUS.md` Bonus 2, and it's the ideal warm-up for Module 5.

THE AGENTIC MOMENT (RECAP)

Guiding question (reflection): *Credential access and exfil left no open thread — nothing pointed at them. What made them findable, and why could the agent only do this now and not in Module 1?*

Manual dead-end: credential access (A9) is buried in normal file activity and DNS exfil (A10) is buried in resolver noise — but the real wall is that *nothing pointed at them.* No open thread. At L0–L2 you stop, believing you're done.

Agentic recovery (L4): `list_coverage_gaps` read your hunts, mapped them to the kill chain, and named the two phases you'd never covered — turning an invisible gap into a to-do list.

MODULE 4 — BONUS / STRETCH HUNTS

> For hunters who finish H-0008 early. Tier 3. See `seeding/TIER_MANIFEST.md`.
> This is also the ideal warm-up for Module 5 — the muscles are identical.

BONUS 1 — DECODE THE EXFIL VOLUME

Guiding question (interpretation): *You confirmed the DNS-tunnel channel exists. Estimate how much data left — what method gets you a bytes-per-query figure, and why is the method more important than the exact number?*

BONUS 2 — ASK THE AGENT TO CONNECT THE WHOLE CAMPAIGN

Guiding question (reflection): *With all five hunts on disk, ask the agent for the single campaign narrative. What story should it stitch, and why is this a preview of the Module 5 payoff?*

```
Use find_similar_hunts across my whole corpus and tell me the single campaign narrative that links them.
```

BONUS 3 — PRE-STAGE FOR MODULE 5

Guiding question (decision): *Run coverage and screenshot it. Which two hunt structures should you know cold before you pivot to the Module 5 campaign — and why those two?*

```
athf hunt coverage --detailed
```

MODULE 5 — LAB: SAME ACTOR, NEW TTP

> **Level coverage: L0–L4 (your choice).** Budget: 15 min. Artifacts: M5-A1 (T1), M5-A3 (T3 ★) in budget; M5-A2 (T2), M5-A4 (T2) as stretch. Deliverable: **H-0009**.
> **Splunk time range: All time.** The telemetry is fixed at 2026-10-23 → 2026-10-30 and does not move with the calendar; a relative window (Last 24 hours, Last 7 days) returns nothing.
> No hand-holding by design. A second campaign — CI-runner telemetry from the same actor — is sitting in `index=thrunt`. Your Campaign-1 hunt repo is your head start. **Start the clock when you pause the video.**

THE RULES

- **No prescribed sequence.** Pick your approach across all five levels. The suggestions below are a fast route, not a script.
- **Reuse aggressively — just ask your repo.** Talk to your corpus in plain English: *"Have I seen this actor before?"*, *"Does this runner beacon match anything I've already hunted?"*, *"Which of my Campaign-1 hunts apply to a CI runner?"* (`athf similar` / `find_similar_hunts` are the explicit fallback.)
- **One hunt minimum (H-0009).** More if you're fast. Cite the earlier hunt you reused.
- **Stuck? Use the answer key.** It ships with the workbook for this module on purpose — reading the hint beats staring. Bring what stumped you to the live debrief.

Guiding question (orientation): *Same actor, new battlefield. Before you pivot to the new campaign — what did CRON GOBLIN do in Campaign 1 that they'd likely do again, and which of your five documented hunts is your best head start?*

FAST ROUTE (15 MIN)

1. The obvious front door — malicious CI build step (M5-A1, Tier 1) (3 min):

Guiding question (interpretation): *A CI runner executes build steps all day. What makes one step malicious rather than a normal pipeline job?*

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
(properties.command_line="*curl*|*sh*" OR properties.command_line="*node -e*")
| table _time hostname principal properties.command_line
```

2. ★ The euphoric-surprise link — same infrastructure (M5-A3, Tier 3) (6 min):

Guiding question (decision): *Ask your repo whether this new campaign shares infrastructure with anything you've already hunted. Which prior hunt do you pivot off, and what would "same actor" proof actually look like in the data?*

Ask your repo *"Does this new campaign share infrastructure with anything I've already hunted?"* — the agent pivots off your Campaign-1 C2 hunt (H-0007). (Fallback: `athf similar --hunt H-0007`.) The connective tissue is the **same C2 domain and hosting IP** you confirmed in H-0007:

```
index=thrunt sourcetype="bro:ssl:json" server_name="cdn-packmule-telemetry.net"
| stats count values(id.resp_h) as c2_ip by server_name
```

Same domain, same `203.0.113.44` — present in *both* campaigns' data. Now attribute the client: egress runs through the forward proxy, so Zeek `ssl` shows the proxy (`10.60.20.30`) as the origin — the **proxy log** carries the true client:

```
index=thrunt sourcetype="apache:access:json" source="*proxy_access.log"
server="cdn-packmule-telemetry.net"
| stats count values(user) as users by client
```

3. Write H-0009 (4 min) — do this even if you only found M5-A1:

Guiding question (decision): *What single line in this hunt record is the entire point of the workshop — the thing that made Campaign 2 fast?*

`bash commands/m5-hunt-new-H0009.sh` or `cp templates/H-0009.md hunts/`, then fill in the runner. For reference:

```
athf hunt new \
  --title "CRON GOBLIN CI-runner poisoning (same-actor pivot from packmule campaign)" \
  --technique T1195.003 \
  --tactic initial-access --tactic command-and-control --tactic exfiltration \
  --platform linux \
  --data-source edr --data-source network \
  --actor "CRON GOBLIN" \
  --behavior "malicious injected CI build step steals runner cloud creds, beacons to Campaign-1 C2 infra, exfils source to attacker cloud storage" \
  --location "<the poisoned CI runner you found in step 1/2>; shared C2 cdn-packmule-telemetry.net" \
  --evidence "eCAR curl|sh + IMDS access; Zeek ssl shared domain across workstation + runner" \
  --threat-context "same actor/infra as H-0005/H-0007 C2 hunts; linked via athf similar. New TTP: CI pipeline poisoning." \
  --hunter "<your-name>" \
  --non-interactive
```

4. Verify (2 min):

```
athf hunt list          # 6 hunts — you leave with the full set
athf hunt coverage       # two campaigns, one actor, mapped
```

Resume the video for the debrief and wrap.

STRETCH (OUTSIDE THE 15 MIN — FINISH THESE BEFORE THE LIVE DEBRIEF IF YOU CAN)

S1. Reuse your credential-access hunt (M5-A2, Tier 2): H-0008 taught you the cred-theft shape. Apply it to the runner:

```
index=thrunt sourcetype="evidenceforge:ecar:json" object=PROCESS action=CREATE
(properties.command_line="*169.254.169.254*" OR properties.command_line="*CI_JOB_TOKEN*")
| table _time hostname principal properties.command_line
```

S2. Close it — exfil to cloud storage (M5-A4, Tier 2):

Guiding question (interpretation): *You'll see traffic to S3. One is exfil, one is a legitimate release. What distinguishes them — and why is the red herring here on purpose?*

```
index=thrunt sourcetype="bro:ssl:json" server_name="*.s3.us-west-2.amazonaws.com"
| stats count by id.orig_h server_name
```

THE AGENTIC MOMENT (RECAP)

Guiding question (reflection): *Two hunters walk into Module 5 — one with a documented repo, one without. Trace exactly where their 15 minutes diverge, and why that gap is the whole workshop.*

Manual dead-end: fresh CI telemetry, an unfamiliar battlefield, 15 minutes. The ad-hoc hunter starts cold and never connects the runner's beacon to the earlier campaign — they see two unrelated events.

Agentic recovery (full loop): the hunter with a documented repo runs `athf similar` / `find_similar_hunts` against their Campaign-1 C2 hunt, and the shared infrastructure (M5-A3) links the two campaigns instantly.

MODULE 5 — BONUS / STRETCH HUNTS

> For the fastest hunters, if you finish the M5 fast route and both stretch hunts.
> See `seeding/TIER_MANIFEST.md`.

BONUS 1 — FULL TWO-CAMPAIGN TIMELINE

Guiding question (decision): *Build one timeline spanning both campaigns. What anchors the two together into a single story a threat-intel team would ship?*

BONUS 2 — WRITE THE DETECTION HAND-OFF NOTE

Guiding question (decision): *You've hunted this. What would make it a durable detection? Name the behavior stable enough to alert on — and the ones that aren't.*

BONUS 3 — BREAK THE LINKAGE ON PURPOSE

Guiding question (reflection): *If CRON GOBLIN had rotated to a fresh C2 domain for the CI campaign, would your Campaign-1 repo still have helped? What behavior — not indicator — would still transfer?*